



Technical Design: Multi-Source Candidate Data Transformer


Name: K Soveet Kumar Prusty | **Email:** soveet.prusty@gmail.com

1. Pipeline / Step Breakdown

The system runs via a deterministic pipeline orchestrator:







1. **Ingest & Extract (Adapters)**: Reads the structured CSV using csv-parse and unstructured PDFs using a layered fallback mechanism (pdf-parse -> stream deflation -> optional lazy OCR).
2. **Cluster & Group (Entity Resolution)**: Clusters all records belonging to the same candidate using cryptographic SHA-256 fingerprinting of normalized emails and phone numbers.
3. **Merge & Resolve (Canonical Engine)**: Deterministically merges arrays of raw records into a single CanonicalProfile. It tracks the origin of every field via a provenance array and calculates a transparent overall_confidence score based on data density and source trust.
4. **Project (Output Formatting)**: A decoupled schema projector mutates the rich internal CanonicalProfile into the exact structure requested by the end-user (e.g., dropping fields, renaming keys) based on a runtime JSON configuration.

2. Unstructured Data (PDF) Extraction

Extracting semantic meaning from visually formatted PDFs is inherently fragile. To maximize resilience without dragging in heavy, brittle dependencies, the adapter implements a three-tier fallback strategy:

1. **Tier 1 (pdf-parse)**: Attempts to extract native text streams.
2. **Tier 2 (Stream Deflation)**: If pdf-parse crashes, falls back to manually inflating PDF objects via zlib and stripping layout markers.
3. **Tier 3 (Lazy Platform-Aware OCR)**: For image-only PDFs, the pipeline lazily requires OCR engines at runtime rather than declaring them as hard dependencies (which would risk failing the reviewer's npm install on machines lacking libcairo/Ghostscript). If available, it uses macOS sips for native rasterization, or pdf2pic on Linux. If the optional dependencies are absent, it safely catches the MODULE_NOT_FOUND error, logs a warning, and gracefully degrades to an empty profile instead of crashing.

Note: Kerning/Spacing artifacts (e.g., "Berl in") caused by PDF layout engines are patched using a generalized commonKerningFixes dictionary during pre-processing rather than hardcoding isolated regexes.

3. Conflict Resolution Strategy

When multiple sources provide different values for the same field (e.g., CSV says "5.5 years", PDF implies "6.4 years"), we resolve conflicts using:

- **Trust Hierarchy**: Explicitly defined source weights (e.g., csv = 1.0, pdf = 0.75, github = 0.85). The system defaults to the highest-trust source.
- **Data Density Thresholding**: We prefer non-null arrays/objects. If a high-trust source has null for skills, but a low-trust source has data, we keep the low-trust data but annotate the provenance.
- **Deterministic Array Merging**: For arrays like experience or skills, we deduplicate items based on normalized string similarity to avoid redundant entries.

4. Alternate Data Source (GitHub)

I integrated the GitHub REST API (`/users/{username}`) as a supplementary unstructured source.

- **Fetching:** Handled via standard fetch with rate-limit backoff logic.
- **Merging:** GitHub profiles are treated as distinct candidate entities *unless* their public email perfectly matches an existing candidate's normalized contact key. This prevents polluting a strong CSV candidate with unrelated GitHub data if the username mapping is ambiguous.

5. Extensibility (Web UI)

While the core requirement was a robust CLI, I additionally built a full-stack Web Application (Express API + Vite/React frontend) to demonstrate how the pipeline can be exposed as a microservice. The UI allows drag-and-drop ingestion and visualizes the generated CanonicalProfile and provenance metrics in real-time.